A more efficient algorithm for appending data

James Ko

November 14, 2017

Abstract

This paper introduces **growth arrays**, array-like data structures that are designed for appending elements. When the number of items is not known beforehand but is expected to be large, growth arrays are more efficient than dynamic arrays by a constant factor. They support all operations dynamic arrays do, such as random access, iteration, and insertion or deletion at an index. For operations other than appending, however, they perform the same as or slightly worse than dynamic arrays.

1 Introduction

In imperative languages, the dynamic array is the most common data structure used by programs. People often want to append multiple items to a collection and then iterate it, which dynamic arrays make simple and efficient. But can it be said they have the *most* efficient algorithm for this pattern?

In this paper, I introduce an alternative data structure to the dynamic array, called the **growth array**. It is more efficient than the dynamic array at appending large numbers of items. This is due to how it 'grows' once it cannot fit more items in its buffer.

When dynamic arrays run out of space, they allocate a new buffer, copy the old one's contents into it, then throw the old one away. Growth arrays are less wasteful in this scenario. Instead of discarding the filled buffer, they simply archive it. The new buffer they allocate represents a continuation of the items from the old buffer. For example, if the old buffer contained items $0 - -9$, the new buffer would contain items 10 and beyond. Because of this technique, growth arrays allocate less memory to store the same number of items, and do not need to copy the contents of old buffers.

Growth arrays have caveats, however. Operations other than appending take the same time or longer than for dynamic arrays. In particular, random access is very costly for growth arrays. They are also not contiguous in memory, which gives them poorer locality than dynamic arrays, and prevents them from being passed to external code that accepts contiguous buffers.

It is worth mentioning that if the size of the data is known in advance, both dynamic and growth arrays are unnecessary. One could simply allocate a raw array with the known size, and append items to it just as quickly. Thus, growth arrays are only beneficial for cases where the amount of data to be appended is unknown, but is expected to be large.

2 External Definitions

I assume that the following functions are defined by the runtime. Thus, I will use them in my algorithms without providing definitions for them.

Algorithm 1 Runtime-defined functions

  ▷ Copies *len* items from *source* to *dest*

  Array_copy(source, dest, len)

  ▷ Returns a new array with length *len*

  New_array(len)

  ▷ Returns a new, empty dynamic array

  New_dynamic_array()

I will also use the following constants in my algorithms. Specific values for these constants must be supplied by the person using the algorithms.

initial capacity Denoted by c_0; the capacity of an empty dynamic array. c_0 must be a natural number.

growth factor Denoted by g; the constant ratio of the new capacity to the old one when L grows. g must be greater than 1.

3 Fields and Properties

In later sections, I will implement algorithms for both dynamic and growth arrays. In this section, I will define fields and properties for those algorithms to use. **Fields** are variables associated with an object that may be read from or written to. **Properties** are trivial, constant time methods that do not change state.

Let L be a dynamic array. I will give L the following fields:

- $L.Buf$ – The **buffer**, or raw array, that L stores its items in.

- $L.Size$ – The number of items in L.

As a dynamic array, L is also given the following properties. (**Note:** For pseudocode, I will use := to denote a definition, and = to denote an equivalence check. Functions that return boolean values are suffixed with ?.)

Algorithm 2 Properties (dynamic array)

▷ Returns the capacity of L

$L.Cap := L.Buf.Len$

▷ Returns whether L is full

$L.Full? := L.Size = L.Cap$

The following code should run when L is instantiated:

Algorithm 3 Constructor (dynamic array)

---
procedure $Constructor(L)$

  $L.Buf \leftarrow New\_array(c_0)$

  $L.Size \leftarrow 0$

Now, let L be a growth array. I will give it the following fields:

- $L.Head$ – The **head** of L. It returns the buffer new items are appended to.

- $L.Tail$ – The **tail** of L. (**Important note: The tail is a dynamic array.**) It returns a dynamic array of pointers to buffers that are already filled with items. The tail can be thought of as a two-dimensional array.
  (**Remark:** It may seem strange for a growth array to use the very data structure it is replacing. Lemma 5.4 will show, however, that only $O(\log n)$ many pointers are appended to the tail. Thus, the extra copying and allocations the tail performs is minuscule compared to other work done by the growth array.)

- $L.Size$ – The number of items in L.

- $L.Cap$ – The **capacity** of L. It returns the maximum number of items L can hold before it must grow.

As a growth array, L is also given the following properties:

Algorithm 4 Properties (growth array)

  ▷ Returns whether L is full
  $L.Full? := L.Size = L.Cap$

  ▷ Returns the **head capacity**, or the capacity of $Head$
  $L.Hcap := L.Head.Len$

  ▷ Returns the **head size**, or the size of $Head$
  ▷ **Explanation:** $Cap - Hcap$ is the number of items in $Tail$.
  ▷ $Size - (Cap - Hcap)$ is the number of items not in $Tail$ (meaning, in $Head$).
  $L.Hsize := L.Size - (L.Cap - L.Hcap)$

The following code should run when L is instantiated:

Algorithm 5 Constructor (growth array)
---
  procedure $Constructor(L)$
    $L.Head \leftarrow New\_array(c_0)$
    $L.Tail \leftarrow New\_dynamic\_array()$
    $L.Size \leftarrow 0$
    $L.Cap \leftarrow c_0$

---

4 Asymptotic Notation

4.1 Introduction

Typically, big-O notation is used to analyze the time or space complexity of a function. In order to highlight the benefit of growth arrays, however, I will use the $\sim$ relation to analyze complexity. This is because growth arrays are only better than dynamic ones by a constant factor for certain operations. For example, dynamic arrays might use roughly $2n$ space for appending n items, while growth arrays would use roughly n space. Even though growth arrays are clearly better in this regard, the big-O space complexity for both data structures would be the same, $O(n)$.

My goal is to be able to compare the coefficients of the highest-order terms in both expressions. For example, I would like to take the ratio $2n : n$, see that it is 2, and conclude that dynamic arrays allocate roughly twice as much as growth arrays for large n. However, big-O notation does not support this.

4.2 Definitions

I will mainly use the $\sim$ relation to analyze complexity. It is defined as follows:

$$f \sim g \iff \lim_{n \to \infty} \frac{f}{g} = 1$$

This is read as "f is asymptotic to g" or "f and g are asymptotic." **Note:** f and g are used as shorthand to denote $f(n)$ and $g(n)$, respectively.

Notice that while $2n = O(n)$, $2n \nsim n$. Thus, $\sim$ makes it possible to distinguish between a function that uses n space and one that uses $2n$ space. **Note:** A consequence of this is that bases for logarithms cannot be omitted, like in big-O notation.

I will also introduce two cousins of $\sim$, the relations $\prec$ and $\preceq$. The former is normally defined as $f \prec g \iff \lim_{n \to \infty} \dfrac{f}{g} = 0$. However, for the purpose of this paper their definitions will be:

$$f \prec g \iff \lim_{n \to \infty} \frac{f}{g} < 1$$
$$f \preceq g \iff f \prec g \vee f \sim g$$

I will call these relations **asymptotic inequalities**. The above inequalities are read "f is asymptotically less than g" and "f is asymptotically less than or equal to g," respectively.

I will also use little-o notation in this paper, which is defined as follows:

$$f = o(g) \iff \lim_{n \to \infty} \frac{f}{g} = 0$$

4.3   Properties

I define various properties of the $\sim$ relation here, which proofs in later sections will use. The properties themselves are proved in the appendix.

Note: In the following theorems, variables with the letters f, g, or h denote functions that, for sufficiently large n, are positive and non-decreasing.

The following theorem states that $\sim$ is a valid equivalence relation.

Theorem 4.1. $\sim$ *is reflexive, transitive, and symmetric.*

The following theorem states that lower-order terms may be removed: for example, $(n + \log_2 n) \sim n$. This is a property shared with big-O.

Theorem 4.2. $f + o(f) \sim f$.

Corollary 4.2.1. *If f is unbounded, then $f + c \sim f$ for any constant c.*

The following theorem states that $\sim$ respects addition, multiplication, and division. This is also a property shared with big-O.

Theorem 4.3. *If $f \sim f_0$ and $g \sim g_0$, then*

$$f + g \sim f_0 + g_0$$
$$fg \sim f_0 g_0$$
$$\frac{f}{g} \sim \frac{f_0}{g_0}$$

Corollary 4.3.1. *A $\sim$ relation is preserved when both sides are added, multiplied, or divided by the same positive quantity.*

The following theorem states that asymptotic functions belong to the same big-O class.

Theorem 4.4. *If $f \sim g$ and $g = O(h)$, then $f = O(h)$.*

The following theorem states that asymptotic functions can be interchanged in an asymptotic inequality.

Theorem 4.5. *If $f \sim f_0$ and $f_0 \prec f_1$, then $f \prec f_1$. If $g \sim g_1$ and $g_0 \prec g_1$, then $g_0 \prec g$.*
These statements also hold for $\preceq$.

The following theorem states that asymptotic inequalities can be algebraically manipulated, in the same vein as Corollary 4.3.1.

Theorem 4.6. *A $\prec$ inequality is preserved when both sides are multiplied or divided by the same positive quantity.*
This statement also holds for $\preceq$.

Notice that the above theorem does not include addition. Asymptotic inequalities do not always hold if the same quantity is added to both sides; for example, $1 \prec 2$, but $n + 1 \sim n + 2$. In order for them to hold, certain conditions must be met.

Theorem 4.7. *Suppose $f \prec g$. If $\lim_{n \to \infty} \dfrac{f}{h} > 0$ and $g = O(h)$, then $f + h \prec g + h$.*
This statement also holds for $\preceq$.

5 Common Operations

In this section, I implement commonly used operations for dynamic and growth arrays, then analyze their time complexity. I also analyze the space complexity of operations that allocate memory.

5.1 Appending

Appending is the most common operation done on dynamic arrays. Growth arrays improve the performance of appending in two ways: by allocating less memory, and by reducing the amount of copying.

Dynamic array implementation

I will implement appending for dynamic arrays first. Let L be a dynamic array. The following constants are used by the algorithm:

Algorithm 6 Appending (dynamic array)

1: **procedure** $Append(L, \ item)$
2:      **if** $L.Full?$ **then**
3:          $L.Grow()$
4:      $L.Buf[L.Size] \leftarrow item$
5:      $L.Size \leftarrow L.Size + 1$

6: **procedure** $Grow(L)$
7:      $new\_buf \leftarrow New\_array(g \times L.Size)$
8:      $Array\_copy(L.Buf, \ new\_buf, \ L.Size)$
9:      $L.Buf \leftarrow new\_buf$

Time complexity

It is well-known that appending one item to a dynamic array takes amortized $O(1)$ time. Thus, appending n items takes $O(n)$ time. However, this fact is not very useful for the purpose of comparison— later, it will be shown that growth arrays also take $O(n)$ time to append n items.

It would be helpful if the ratio $\dfrac{t_D(n)}{t_G(n)}$ could be determined, where $t_D(n)$ and $t_G(n)$ denote the average time needed to append n items for dynamic and growth arrays, respectively. This represents how much faster growth arrays are than dynamic ones. However, there is no way to quantify these functions without empirical measurements. Thus, I will approximate this ratio through a different technique: Suppose n elements are appended to an empty collection. How many times is an element stored in an array? I will term the answer to this question the **write cost** of n appends, and denote it $w(n)$.

I will naïvely assume that $t(n) \propto w(n)$, or that the average time for n appends is proportional to their write cost. Then $\dfrac{t_D(n)}{t_G(n)} = \dfrac{w_D(n)}{w_G(n)}$, where w_D and w_G are respectively the write cost functions for dynamic and growth arrays. In this section, I will derive the formula for $w_D(n)$. Because this section concerns dynamic arrays only, I will write it as $w(n)$.

In the code for *Append*, one array store is performed unconditionally, so it is apparent that $w(n) \geq n$ after n appends. However, *Grow* also does some writing, so in order to find a precise formula for $w(n)$, I need to analyze when *Grow* is called. To do this, I use the following lemma:

Lemma 5.1. *Let L be a dynamic array. Let its **capacity sequence**, κ, be the range of values for $L.Cap$ as n items are appended. For $n = 0$, trivially $\kappa = (c_0)$. For $n > 0$,*

$$\kappa = c_0, \; gc_0, \; g^2c_0, \; \ldots \; g^{\max\left(\left\lceil \log_g n - \log_g c_0 \right\rceil, 0\right)} c_0$$

Proof. I use the following properties of dynamic arrays:

1. The capacity of an empty dynamic array is c_0.

2. The capacity of a dynamic array can only grow by g.

3. The capacity is as small as possible. Put formally, if κ_i is the capacity for n items, then $\kappa_i \geq n$ but $n > \kappa_{i-1}$. (By convention, $\kappa_{-1} = 0$.)

Assumption (1) immediately shows $\kappa_0 = c_0$. Assumption (2) shows that if $g^i c_0$ is the current capacity, then $g^{i+1} c_0$ must be the next capacity. By induction, $\kappa = \left(g^i c_0\right)_{i=0}^{\lambda}$ for some whole number λ.

The final value of the sequence, κ_λ, is the capacity needed to store n items. By assumption (3), $\kappa_\lambda \geq n > \kappa_{\lambda-1}$. Consider the case when $n > c_0$: it must be true that $\kappa_\lambda > c_0$, so $\lambda \geq 1$. Since $\lambda - 1 \geq 0$, $\kappa_\lambda = g^\lambda c_0$ and $\kappa_{\lambda-1} = g^{\lambda-1} c_0$. Then

$$g^\lambda c_0 \geq n > g^{\lambda-1} c_0$$
$$g^\lambda \geq \frac{n}{c_0} > g^{\lambda-1}$$
$$\lambda \geq \log_g n - \log_g c_0 > \lambda - 1$$

Since λ is an integer,

$$\lambda = \left\lceil \log_g n - \log_g c_0 \right\rceil$$

Now, consider the case when $n \leq c_0$. By assumption (3), $n > \kappa_{\lambda-1}$. $\lambda - 1$ must then equal -1, since any other value would imply $n > \kappa_{\lambda-1} \geq c_0$. Thus $\lambda = 0$.

It was shown that $\lambda \geq 1 \geq 0$ for the first case, and it can be shown that $\lceil \log_g n - \log_g c_0 \rceil \leq 0$ for the second case. Therefore, a general formula for λ is as follows:

$$\lambda = \max(\lceil \log_g n - \log_g c_0 \rceil, 0)$$

The final term in the sequence is $g^\lambda c_0 = g^{\max(\lceil \log_g n - \log_g c_0 \rceil, 0)} c_0$, completing the proof. $\square$

Corollary 5.1.1. *Let the **growth sequence**, γ, of L be the sizes at which Grow is called when n items are appended to L. Then $\gamma = \kappa \setminus \{\kappa_\lambda\}$.*

Proof. If κ_i exists and $i \geq 1$, then clearly *Grow* must have been called when the size was κ_{i-1}, so $\kappa_{i-1} \in \gamma$. Then γ contains every term in κ except for the last, κ_λ, as the corollary states. $\square$

When *Grow* is called and the current size is γ_i, the algorithm copies γ_i items. Then the total number of items copied when n items are appended is:

$$\sum_i \gamma_i = c_0 + g c_0 + \ldots + g^{\lambda-1} c_0$$

$$= \left(\frac{g^\lambda - 1}{g - 1} \right) c_0$$

Counting the writes made per item by *Append*, an explicit formula for $w(n)$ is as follows:

$$w(n) = n + \left(\frac{g^\lambda - 1}{g - 1} \right) c_0$$

Now, my goal is to approximate $w(n)$ with $\sim$. To make this easier to do, I will asymptotically bound g^λ with respect to n.

Lemma 5.2. $\dfrac{n}{c_0} \preceq g^\lambda \prec \dfrac{gn}{c_0}$. *That is, $\dfrac{n}{c_0} \leq g^\lambda < \dfrac{gn}{c_0}$ for sufficiently large n.*

Proof. It was shown in Lemma 5.1 that if $n > c_0$, $\lambda = \lceil \log_g n - \log_g c_0 \rceil \geq 1$. Now, note that λ may be written as $\lceil \log_g \frac{n}{c_0} \rceil$ for such n. Then,

$$\log_g \frac{n}{c_0} \leq \lambda < \log_g \frac{n}{c_0} + 1$$
$$\frac{n}{c_0} \leq \lambda < \frac{gn}{c_0}$$

as desired. $\square$

I can build on this inequality to receive asymptotic bounds for $w(n)$:

$$\frac{n}{c_0} \preceq g^\lambda \prec \frac{gn}{c_0}$$
$$\frac{n}{c_0} \preceq g^\lambda - 1 \prec \frac{gn}{c_0} \qquad\qquad \text{(Corollary 4.2.1, Theorem 4.5, Corollary ??)}$$
$$\frac{n}{g-1} \preceq \left(\frac{g^\lambda - 1}{g - 1} \right) c_0 \prec \frac{gn}{g-1} \qquad\qquad \text{(Theorem 4.6, Corollary ??)}$$
$$\left(\frac{g}{g-1} \right) n \preceq w(n) \prec \left(\frac{2g-1}{g-1} \right) n \qquad\qquad \text{(Theorem 4.7, Corollary ??)}$$

Space complexity

I wish to find the space allocated when n items are appended to a collection. I call this quantity the **space cost**, denote it $s(n)$, and define it as the total length of buffers allocated by n appends. Now, I derive a formula for $s(n)$.

First, from the definition of $L.Cap$, note that a dynamic array's capacity is the length of the buffer it stores its items in. Then a buffer of length c is allocated at some point if and only if $c \in \kappa$. Then the total length of those buffers is

$$s(n) = \sum_i \kappa_i$$
$$= c_0 + gc_0 + g^2 c_0 + \ldots + g^\lambda c_0$$
$$= \left(\frac{g^{\lambda+1} - 1}{g - 1} \right) c_0$$

Using Lemma 5.2 again, I asymptotically bound $s(n)$:

$$\frac{n}{c_0} \preceq g^\lambda \prec \frac{gn}{c_0}$$

$$\frac{gn}{c_0} \preceq g^{\lambda+1} \prec \frac{g^2 n}{c_0} \qquad \text{(Theorem 4.6, Corollary ??)}$$

$$\frac{gn}{c_0} \preceq g^{\lambda+1} - 1 \prec \frac{g^2 n}{c_0} \qquad \text{(Corollary 4.2.1, Theorem 4.5, Corollary ??)}$$

$$\left(\frac{g}{g-1} \right) n \preceq s(n) \prec \left(\frac{g^2}{g-1} \right) n \qquad \text{(Theorem 4.6, Corollary ??)}$$

Growth array implementation

Figure 1: Appending 30 items to a growth array with $c_0 = 4$, $g = 2$


Figure 2: Appending 16 items to a growth array with $c_0 = 2$, $g = 3$

In this section, L is a growth array. The following algorithm implements appending for growth arrays.

Algorithm 7 Appending (growth array)

 1: **procedure** $Append(L,\ item)$

 2:     **if** $L.Full?$ **then**

 3:        $L.Grow()$

 4:     $L.Head[L.Hsize] \leftarrow item$

 5:     $L.Size \leftarrow L.Size + 1$

 6: **procedure** $Grow(L)$

 7:     $L.Tail.Append(L.Head)$

 8:     **if** $L.Cap = c_0$ **then**

 9:        $new\_hcap \leftarrow (g-1) \times c_0$

10:     **else**

11:        $new\_hcap \leftarrow g \times L.Hcap$

12:     $L.Head \leftarrow New\_array(new\_hcap)$

13:     $L.Cap \leftarrow L.Cap + new\_hcap$

Time complexity

It takes $O(n)$ time to append n items to a growth array. This can be easily seen by noting that every statement in the algorithm has amortized $O(1)$ time complexity (ignoring the call to New_array, since memory allocation is non-deterministic).

Now, I focus on finding a formula for the write cost function, $w_G(n)$, which will be written as $w(n)$ in this section. I claim that Lemma 5.1 still holds, since growth arrays satisfy all properties used by that proof. In particular, although growth arrays use a different growth algorithm than dynamic arrays, the following statement still holds:

Lemma 5.3. *The capacity of a growth array grows by the constant factor g.*

Proof. I prove that the *Grow* algorithm enforces this using induction. I induct on the number of times *Grow* is called, k, showing that for all natural numbers k, *Grow* behaves correctly when called the kth time. I will let c_i and c_f denote the initial/final capacities and h_i and h_f denote the initial/final head capacities for the kth call, respectively.

For $k = 1$, $c_i = c_0$. I wish to show that $c_f = gc_0$. This happens if and only if the next buffer has size $\Delta c = (g-1)c_0$, which the algorithm ensures.

For $k > 1$, by induction $c_i = $ previous $c_f = g^{k-1}c_0$, and $h_i = $ previous $h_f = (g^{k-1} - g^{k-2})c_0$. I wish to show $c_f = g^k c_0$ and $h_f = (g^k - g^{k-1})c_0$. Because $k > 1$, the algorithm will calculate h_f as g times h_i. Then

$$h_f = gh_i = g(g^{k-1} - g^{k-2})c_0 = (g^k - g^{k-1})c_0$$

and

$$c_f = c_i + h_f = g^{k-1}c_0 + (g^k - g^{k-1})c_0 = g^k c_0$$

as desired. $\qquad\square$

Since Lemma 5.3 has been proven, Lemma 5.1 and all results based on it must also hold true for growth arrays. Now, I am ready to find the write cost of *Grow*. Unlike dynamic arrays, *Grow* does not make γ_i writes when the current size is γ_i. In fact, *Grow* does not copy *any* items supplied by the user. Writes are only made when a buffer pointer is appended to the tail, since the tail is a dynamic array.

Let $w_\gamma(n)$ denote the total number of writes made by *Grow*, and let n_τ be the size of the tail. Since Corollary 5.1.1 also holds true for growth arrays, *Grow* is called $|\gamma|$ times. A buffer is appended to the tail each time *Grow* is called. Thus, the tail's size is

$$n_\tau = |\gamma| = |\kappa| - 1 = \lambda$$

Then the formula for $w_\gamma(n)$ is simply $w_D(\lambda)$, since λ pointers are appended to the tail, which is a dynamic array. Finally, adding the writes made by *Append*, the formula for $w(n)$ is

$$w(n) = n + w_D(\lambda)$$

Now, I approximate $w(n)$ using $\sim$. To do this, I will derive the big-O complexity of λ.

Lemma 5.4. $\lambda = O(\log n)$.

Proof. From Lemma 5.1, $\lambda = \max(\lceil \log_g n - \log_g c_0 \rceil, 0)$. For sufficiently large n, $\lambda = \lceil \log_g n - \log_g c_0 \rceil$. Using $\sim$ to analyze λ, I receive

$$\lambda \sim \lceil \log_g n - \log_g c_0 \rceil \sim \left(\log_g n - \log_g c_0 \right) \sim \log_g n$$

(It is trivial to show that $\lceil f(n) \rceil \sim f(n)$ if f is unbounded, since $\lceil f(n) \rceil - f(n)$ is bounded by 1.)

By Theorem 4.4, $O(\lambda) = O(\log_g n) = O(\log n)$ as desired. $\qquad\square$

Now when $w(n)$ is approximated with $\sim$, $w_D(\lambda)$ disappears:

$$w(n) = n + w_D(\lambda) = n + O(\lambda) = n + O(\log n) \sim n$$

The last step is justified by Theorem 4.2.

Space complexity

Unlike dynamic arrays, growth arrays do not throw away buffers. This means that if the current capacity is c, then the total length of buffers allocated to store items is also c. Typically, however, $s(n) > c$. This is because the tail of growth arrays (which is a dynamic array) also allocates buffers to store buffer pointers. The space the tail allocates must be considered in the formula for $s(n)$.

In the previous section, I established that $n_\tau = \lambda$. Since κ_λ is the capacity needed to hold n items, and $s_D(\lambda)$ is the space the tail needs to store λ pointers, I conclude that the formula for $s(n)$ is

$$s(n) = \kappa_\lambda + s_D(\lambda)$$

(s_D denotes the space cost function for dynamic arrays.)

Now, I wish to approximate this using $\sim$. Since $\kappa_\lambda = g^\lambda c_0$, it follows from Lemma 5.2 that $n \preceq \kappa_\lambda \prec gn$. Since n and gn are both $O(n)$, [the squeeze theorem] implies that $\kappa_\lambda = O(n)$.

Now, consider the ratio $\lim_{n \to \infty} \dfrac{s_D(\lambda)}{\kappa_\lambda} = \dfrac{O(\log n)}{O(n)} = O\left(\dfrac{\log n}{n}\right) = 0$. Because it is 0, from Theorem 4.2 one can conclude

$$s(n) = \kappa_\lambda + s_D(\lambda) \sim \kappa_\lambda$$

Substituting $s(n)$ for κ_λ in the previous inequality,

$$n \preceq s(n) \prec gn$$

Time complexity comparison

In this section, w_D and w_G will denote the write cost functions for dynamic and growth arrays, respectively. I derived earlier that

$$\left(\frac{g}{g-1}\right) n \preceq w_D(n) \prec \left(\frac{2g-1}{g-1}\right) n$$

From Theorem 4.6 and Corollary **??**, I may divide all sides of this inequality by n to receive

$$\frac{g}{g-1} \preceq \frac{w_D(n)}{n} \prec \frac{2g-1}{g-1}$$

It was shown earlier that $w_G(n) \sim n$. From Theorem 4.3, $\dfrac{w_D(n)}{n} \sim \dfrac{w_D(n)}{w_G(n)}$. From Theorem 4.5 and Corollary **??**, it follows that

$$\frac{g}{g-1} \preceq \frac{w_D(n)}{w_G(n)} \prec \frac{2g-1}{g-1}$$

Space complexity comparison

Let $s_D(n)$ and $s_G(n)$ be the space cost functions for dynamic and growth arrays, respectively. Recall that

$$s_D(n) = \left(\frac{g^{\lambda+1} - 1}{g - 1}\right) c_0$$

and that $s_G(n) \sim \kappa_\lambda = g^\lambda c_0$. Taking the ratio of $s_D(n)$ to $s_G(n)$, I receive

$$\frac{s_D(n)}{s_G(n)} \sim \frac{1}{g^\lambda c_0} \cdot \left(\frac{g^{\lambda+1} - 1}{g - 1}\right) c_0 \qquad \text{(Theorem 4.3)}$$
$$= \frac{g - 1/g^\lambda}{g - 1}$$
$$\sim \frac{g}{g - 1} \qquad \text{(Theorem 4.2)}$$

5.2   Indexing

Indexing, or accessing the ith element of a collection, is another common dynamic array operation. Indexing a growth array involves more instructions than indexing a dynamic array. Depending on how it is implemented, indexing for growth arrays runs in either constant or logarithmic time.

Dynamic array implementation

The following algorithm implements *Get_item* for dynamic arrays. (A *Set_item* function can be implemented in a similar fashion.)

Note: My algorithms assume all arguments passed to them are valid. Thus, arguments are never checked.

Note: Although the function is called *Get_item*, the syntax other algorithms use to call it is $L[index]$.

Algorithm 8 Random access (dynamic array)

1: **function** $Get\_item(L,\ index)$
2:      **return** $L.Buf[index]$

This algorithm runs in $O(1)$ time. It also has the benefit of using very few instructions, which is a benefit that will not be shared with growth arrays.

Growth array implementation

Indexing a growth array is slower and more complex than indexing a dynamic array. However, this does not mean that accessing data of a growth array is necessarily slower than accessing that of a dynamic array. Refer to Sections 5.3 and 5.4 for methods to quickly access growth arrays' data.

Despite growth arrays' relatively poor indexing performance, I include this section for two reasons. 1) Dynamic arrays are indexed frequently. In order for growth arrays to be a viable replacement for them, it should be possible to index growth arrays too. 2) I wish to show that it is possible to index growth arrays in constant time.

$O(\log n)$ **implementation**

The first algorithm I will demonstrate is a naïve implementation that runs in logarithmic time:

Algorithm 9 Random access (growth array), logarithmic time

---

    function $Get\_item(L,\ index)$

        $i \leftarrow index$

        for buf **in** $L.Tail$ **do**

            if $i < buf.Len$ **then**

                return $buf[i]$

            $i \leftarrow i - buf.Len$

        return $L.Head[i]$

I know that this algorithm runs in $O(\log n)$ time since it was shown earlier that the size of the tail, n_τ, equals λ, and (from Lemma 5.4) that $\lambda = O(\log n)$. Aside from iteration over the tail, all statements run in constant time, so the time complexity is precisely $O(\log n)$.

I present this algorithm alongside the constant time one because it is easier to understand, and the latter is often slower if special circumstances are not met.

$O(1)$ **implementation**

In this section, I will present a constant time algorithm for indexing a growth array. Before I do so, however, I must establish a mathematical justification for it.

Suppose I want to get the ith element of a growth array, where i is zero-based. I assume that i is a valid index, or that $0 \leq i < L.Size$. In order to locate the desired element, I must find two things: the buffer that holds the item, and the index of the item within that buffer.

Now, consider that every buffer except the head is stored inside the tail. Thus, such buffers can be uniquely identified by their index in the tail. I will refer to this quantity as the **buffer index** and denote it i_B. I wish to assign the head a buffer index as well, so that each buffer has a unique ID. Since the head follows the tail's last buffer, which has $i_B = n_\tau - 1$, I will let the head's i_B be n_τ.

I define a helper function, Get_buf, that gets the buffer associated with a given i_B. i_B is assumed to be valid; that is, $0 \leq i_B \leq n_\tau$.

Algorithm 10 Helper function
---
function $Get\_buf(L,\ i_B)$
    if $i_B < L.Tail.Size$ **then**
        return $L.Tail[i_B]$
    else
        return $L.Head$

I will call the index of the desired element within the buffer the **element index**, and denote it i_E.

If I define a helper function *Decompose* that returns the i_B and i_E associated with i in an ordered pair, then *Get_item* may be written as follows:

Algorithm 11 Random access (growth array), constant time
---
function $Get\_item(L,\ index)$
    $(i_B,\ i_E) \leftarrow Decompose(index)$
    return $L.Get\_buf(i_B)[i_E]$

Now, the task is to find formulae for i_B and i_E in terms of i, in order to implement *Decompose*.

Lemma 5.5. *The formulae for i_B and i_E are*

$$i_B = \lambda|_{n=i+1}$$
$$i_E = i - \gamma_{i_B-1}$$

(By convention, $\gamma_{-1} = 0$.)

Proof. Appending new items does not change the index of an item that is already in the list. Thus, finding the element at index i is equivalent to finding the last element when $n = i + 1$.

The last element always resides in the head, so $i_B = n_\tau|_{n=i+1}$. As shown earlier, $n_\tau = \lambda$, so $i_B = \lambda|_{n=i+1}$.

To determine i_E, consider the identity:

$$n = \#\text{ of items in tail} + \#\text{ of items in head}$$

First, I will determine the number of items in the tail. In the case where $i_B = 0$, then because $n_\tau = i_B$, the tail contains 0 buffers and thus 0 items (which equals γ_{-1}). If $i_B > 0$, the number of items in the tail is the last size at which *Grow* was called, or the last term of γ. This quantity is $\gamma_{\lambda-1}$, or γ_{i_B-1}.

The desired item is the growth array's last element, which implies it is also the head's last element. Thus i_E is the head's last valid index, so the head's size is $i_E + 1$. Finally, from the premise, $n = i + 1$. Substituting all values into the above identity, I receive

$$i + 1 = \gamma_{i_B-1} + (i_E + 1)$$
$$i_E = i - \gamma_{i_B-1}$$

completing the proof. $\square$

Using the formulae for λ and γ_i, I now implement the *Decompose* function:

Algorithm 12 Helper function

function $Decompose(index)$
    $i_B \leftarrow \max(\lceil \log_g(index + 1) - \log_g c_0 \rceil, 0)$
    if $i_B > 0$ **then**
        $i_E \leftarrow index - g^{i_B - 1} \times c_0$
    else
        $i_E \leftarrow index$
    return (i_B, i_E)

Clearly, *Decompose* runs in constant time. Despite that, it appears to be quite expensive: normally, logarithms and exponentiation require use of floating-point instructions, which are slower than integer-based instructions. However, in the special case where $g = 2$ and $c_0 = 2^\varepsilon$ for some constant whole number ε, I claim that i_B and i_E can be computed without use of floating-point instructions.

To see this, first note that i_B becomes $\max(\lceil \log_2(i + 1) \rceil - \varepsilon, 0)$. There is a constant time algorithm to compute $\lceil \log_2 k \rceil$ without using floating-point for any $k \in \mathbb{N}$, which I will not discuss because it involves concepts outside the scope of this paper. (An implementation of the algorithm, however, may be found in the links at Section 6.) Equipped with such an algorithm, it is easy to see that the whole expression can be computed without use of floating-point.

In calculating i_E, the only potential use of floating-point instructions comes from $g^{i_B - 1}$. When $g = 2$, however, this can be calculated with a simple bit shift. Thus, it is not necessary to use floating-point instructions to calculate either i_B or i_E.

5.3 Iterating

Informally, to **iterate** a collection is to loop over each of its elements. The syntax "**for** *item* **in** *collection*" implicitly uses the iteration algorithm of *collection*.

It is very common to iterate a dynamic array once items are appended to it. In this section, I will show that growth arrays perform comparably to dynamic arrays for iteration.

Dynamic array implementation

The algorithm for iterating dynamic arrays is simple and familiar.

Algorithm 13 "**for** *item* **in** L" algorithm (dynamic array)

for $i = 0$, $i < L.Size$, $i \leftarrow i + 1$ **do**
    $item \leftarrow L[i]$
    do something with $item$

Growth array implementation

The iteration algorithm for growth arrays is not quite as simple. As shown in the previous section, growth arrays' random access algorithm is expensive; thus, it is best avoided. The following algorithm implements iteration without using the random access algorithm.

Algorithm 14 "**for** *item* **in** *L*" algorithm (growth array)

---

   for *buf* **in** *L.Tail* **do**

      for *item* **in** *buf* **do**

         do something with *item*

   for $i = 0$, $i < L.Hsize$, $i \leftarrow i + 1$ **do**

      $item \leftarrow L.Head[i]$

      do something with *item*

5.4 Copying to an array

Users often want to convert dynamic arrays to raw arrays. There are multiple reasons for this: 1) Raw arrays hold on to exactly n memory to store n elements. However, dynamic and growth arrays typically use more than n memory, so they need not grow every time an item is appended. 2) Sometimes, functions in third-party code only accept raw arrays.

For growth arrays, the case is even more compelling: 3) Since they are not contiguous as a whole, they have worse data locality (even if the number of discontinuities is $O(\log n)$). 4) Raw arrays can be indexed much faster than growth arrays.

In this section, I will show that growth arrays can be converted to raw arrays not much slower than dynamic arrays.

Dynamic array implementation

It is straightforward to convert a dynamic array to a raw array:

Algorithm 15 Converting to raw array (dynamic array)

---

   function $To\_raw\_array(L)$

      $raw\_array \leftarrow New\_array(L.Size)$

      $Array\_copy(L.Buf,\ raw\_array,\ L.Size)$

      return raw_array

Since n elements must be copied, the time complexity of this function is $O(n)$.

Growth array implementation

The growth array algorithm is slightly slower than the dynamic array one. It needs to loop over the tail and call *Array_copy* multiple times, since not all elements are contiguous. However, n elements are still copied in total, so the time complexity is still $O(n)$.

Algorithm 16 Converting to raw array (growth array)

---

   function $To\_raw\_array(L)$

      $raw\_array \leftarrow New\_array(L.Size)$

      for *buf* **in** *L.Tail* **do**

         $Array\_copy(buf,\ raw\_array,\ buf.Len)$

      $Array\_copy(L.Head,\ raw\_array,\ L.Hsize)$

      return raw_array

5.5 Other Operations

Growth arrays support other dynamic array operations, such as insertion/deletion at an index, binary search, sorting, etc. I will not give algorithms for these operations in this paper, since they are not as common as the four detailed above. However, I will briefly describe how insertion and deletion can be implemented.

Suppose the user wants to insert an item at index i. If $i = n$, then simply append the item. Otherwise $(0 \leq i < n)$, capture $L[n-1]$ in a local variable. Shift all elements with index $\geq i$ forward by one, so that their index increases by one. Append the original value of $L[n-1]$. Finally, set index i to the provided item.

Suppose the user wants to delete the item at index i. Shift all elements with index $> i$ backwards by one, so that their index decreases by one. Decrement $L.Size$ and $L.Hsize$. If this causes $L.Hsize$ to become 0 and $n_\tau > 0$, then discard the head and replace it with the last buffer pointed to by the tail. Remove this buffer from the tail.

6 Implementations

7 Closing Remarks

Appendix

A   Proofs of $\sim$ Properties

Note: In the following theorems and their proofs, variables with the letters f, g, or h denote functions that, for sufficiently large n, are positive and non-decreasing.

A.1 $\sim$ is an Equivalence Relation

Theorem 4.1. $\sim$ *is reflexive, transitive, and symmetric.*

Proof. Clearly, $\lim_{n\to\infty} \dfrac{f}{f} = 1$, so $f \sim f$. Thus, $\sim$ is reflexive.

Suppose $f \sim g$. Then $\lim_{n\to\infty} \dfrac{f}{g} = 1$. Since both $\lim_{n\to\infty} 1$ and $\lim_{n\to\infty} \dfrac{f}{g}$ exist and the latter is nonzero, $\dfrac{\lim_{n\to\infty} 1}{\lim_{n\to\infty}(f/g)} = \lim_{n\to\infty} \dfrac{g}{f}$. Since the left-hand side evaluates to 1, $g \sim f$. Thus, $\sim$ is symmetric.

Suppose $f \sim g$ and $g \sim h$. By definition, $\lim_{n\to\infty} \dfrac{f}{g} = 1$ and $\lim_{n\to\infty} \dfrac{g}{h} = 1$. Since both limits exist, their product is $\lim_{n\to\infty} \left(\dfrac{f}{g} \cdot \dfrac{g}{h} \right) = \lim_{n\to\infty} \dfrac{f}{h}$. Since this product is 1, $f \sim h$. Thus, $\sim$ is transitive. $\quad\square$

A.2 $\sim$ Removes Lower-Order Terms

Theorem 4.2. $f + o(f) \sim f$.

Proof. Let $g = o(f)$. By definition, $\lim_{n\to\infty} \dfrac{g}{f} = 0$. Since $\lim_{n\to\infty} \dfrac{f}{f}$ and $\lim_{n\to\infty} \dfrac{g}{f}$ both exist,

$$\lim_{n\to\infty} \frac{f+g}{f} = \lim_{n\to\infty} \frac{f}{f} + \lim_{n\to\infty} \frac{g}{f} = 1 + 0 = 1$$

It follows that $f + g = f + o(f) \sim f$. $\quad\square$

A.2.1 $f + c \sim f$ for Unbounded f

Corollary 4.2.1. *If f is unbounded, then $f + c \sim f$ for any constant c.*

Proof. If f is unbounded, then $c = o(f)$. Applying Theorem 4.2, $f + c \sim f$. $\quad\square$

A.3 $\sim$ Respects $+$, $\times$, and $\div$

Theorem 4.3. *If $f \sim f_0$ and $g \sim g_0$, then*

$$f + g \sim f_0 + g_0$$
$$fg \sim f_0 g_0$$
$$\frac{f}{g} \sim \frac{f_0}{g_0}$$

Proof. Let $q = \dfrac{f+g}{f_0 + g_0}$, and let $d = g_0 + g_0^2/f_0$. q may be expressed as

$$\begin{aligned}
q &= \frac{(g_0/f_0)(f+g)}{(g_0/f_0)(f_0 + g_0)} \\
&= \frac{g_0(f/f_0) + g_0 g/f_0}{d} \\
&= \frac{g_0(f/f_0) + (g_0^2/f_0)(g/g_0)}{d} \\
&= \frac{g_0}{d} \cdot \frac{f}{f_0} + \frac{g_0^2/f_0}{d} \cdot \frac{g}{g_0}
\end{aligned}$$

Now, consider that

$$\begin{aligned}
q - 1 = q - \frac{d}{d} &= q - \frac{g_0}{d} - \frac{g_0^2/f_0}{d} \\
&= \frac{g_0}{d} \cdot \frac{f}{f_0} + \frac{g_0^2/f_0}{d} \cdot \frac{g}{g_0} - \frac{g_0}{d} - \frac{g_0^2/f_0}{d} \\
&= \frac{g_0}{d} \cdot \left(\frac{f}{f_0} - 1 \right) + \frac{g_0^2/f_0}{d} \cdot \left(\frac{g}{g_0} - 1 \right) \\
\lim_{n \to \infty} (q - 1) &= \lim_{n \to \infty} \left(\frac{g_0}{d} \cdot \left(\frac{f}{f_0} - 1 \right) + \frac{g_0^2/f_0}{d} \cdot \left(\frac{g}{g_0} - 1 \right) \right) \\
&= \lim_{n \to \infty} \left(\frac{g_0}{d} \cdot \left(\frac{f}{f_0} - 1 \right) \right) + \lim_{n \to \infty} \left(\frac{g_0^2/f_0}{d} \cdot \left(\frac{g}{g_0} - 1 \right) \right)
\end{aligned}$$

Since g_0 and g_0^2/f_0 sum to d and all functions are positive, $\dfrac{g_0}{d}$ and $\dfrac{g_0^2/f_0}{d}$ are bounded between 0 and 1 as $n \to \infty$. From the given, $\lim_{n \to \infty} \left(\dfrac{f}{f_0} - 1 \right)$ and $\lim_{n \to \infty} \left(\dfrac{g}{g_0} - 1 \right)$ both exist and equal 0. The limit of a bounded expression times one approaching 0 is 0; thus, both limits on the right-hand side are 0.

Substituting, I receive $\lim_{n \to \infty}(q - 1) = 0$, so $\lim_{n \to \infty} q = 1$. Since q is defined as $\dfrac{f+g}{f_0 + g_0}$, this shows that $f + g \sim f_0 + g_0$, proving the theorem statement for addition.

The statement is proven much more easily for multiplication. Since both $\lim_{n \to \infty} \dfrac{f}{f_0}$ and $\lim_{n \to \infty} \dfrac{g}{g_0}$ exist,

$$\lim_{n \to \infty} \frac{fg}{f_0 g_0} = \left(\lim_{n \to \infty} \frac{f}{f_0} \right) \left(\lim_{n \to \infty} \frac{g}{g_0} \right) = 1 \cdot 1 = 1$$

It follows that $fg \sim f_0 g_0$. If the limit for g is flipped before multiplying, the following results:

$$\lim_{n \to \infty} \frac{f/g}{f_0/g_0} = 1$$

This implies that $\dfrac{f}{g} \sim \dfrac{f_0}{g_0}$. $\qquad \square$

A.3.1 $\sim$ Equations can be Algebraically Manipulated

Corollary 4.3.1. *A $\sim$ relation is preserved when both sides are added, multiplied, or divided by the same positive quantity.*

Proof. This immediately follows from Theorem 4.3 by taking $g_0 = g$. $\qquad \square$

A.4 Asymptotic Functions Have the Same Big-O Class

Theorem 4.4. *If $f \sim g$ and $g = O(h)$, then $f = O(h)$.*

Proof. Let $r = \dfrac{f}{g}$. Since $\lim_{n \to \infty} r = 1$, $r \leq 2$ for sufficiently large n. By definition, $\exists c : g \leq ch$ for sufficiently large n, where c is a positive constant. Then, $f = rg \leq rch \leq 2ch$ for sufficiently large n. Thus, $f = O(h)$. $\qquad\square$

A.5 Asymptotic Functions may be Interchanged in $\prec$ and $\preceq$ Inequalities

Theorem 4.5. *If $f \sim f_0$ and $f_0 \prec f_1$, then $f \prec f_1$. If $g \sim g_1$ and $g_0 \prec g_1$, then $g_0 \prec g$.*
These statements also hold for $\preceq$.

Proof. By definition, $\lim_{n \to \infty} \dfrac{f}{f_0} = 1$ and $\lim_{n \to \infty} \dfrac{f_0}{f_1} < 1$. Since both limits exist,

$$\lim_{n \to \infty} \frac{f}{f_1} = \left(\lim_{n \to \infty} \frac{f}{f_0} \right) \left(\lim_{n \to \infty} \frac{f_0}{f_1} \right) < 1 \cdot 1 = 1$$

It follows that $f \prec f_1$.

The second statement may be proved in a similar manner. Since $\sim$ is symmetric, $g_1 \sim g$. Applying the same argument as before,

$$\lim_{n \to \infty} \frac{g_0}{g} = \left(\lim_{n \to \infty} \frac{g_0}{g_1} \right) \left(\lim_{n \to \infty} \frac{g_1}{g} \right) < 1 \cdot 1 = 1$$

Thus, $g_0 \prec g$.

From these results and the transitivity of $\sim$, it can trivially be shown that the same statements are true when $\prec$ is replaced with $\preceq$. $\qquad\square$

A.6 $\prec$ and $\preceq$ Inequalities can be Algebraically Manipulated

Theorem 4.6. *A $\prec$ inequality is preserved when both sides are multiplied or divided by the same positive quantity.*
This statement also holds for $\preceq$.

Proof. Suppose that $f \prec g$. By definition, $\lim_{n \to \infty} \dfrac{f}{g} < 1$. For any positive function h, $\lim_{n \to \infty} \dfrac{fh}{gh} < 1$ and $\lim_{n \to \infty} \dfrac{f/h}{g/h} < 1$, so respectively $fh \prec gh$ and $\dfrac{f}{h} \prec \dfrac{g}{h}$.

From these results and Corollary 4.3.1, it can trivially be shown that the same statement is true when $\prec$ is replaced with $\preceq$. $\qquad\square$

A.7 Same-Order Terms may be Added to Both Sides of $\prec$ and $\preceq$ Inequalities

Theorem 4.7. *Suppose $f \prec g$. If $\lim_{n \to \infty} \dfrac{f}{h} > 0$ and $g = O(h)$, then $f + h \prec g + h$.*
This statement also holds for $\preceq$.

Proof. Assume that n is arbitrarily large. Since $\dfrac{f}{h}$ has a positive limit, it is bounded below by some positive constant c_1. Rearranging, $c_1 h \leq f$. From $f \prec g$, it follows that $f < g$ for sufficiently large n. Since $g = O(h)$, $\exists c_2 : g \leq c_2 h$. Combining these inequalities, I receive $c_1 h \leq f < g \leq c_2 h$. From this, it is clear that $c_1 < c_2$.

Now, consider that

$$\lim_{n\to\infty} \frac{f+h}{g+h} \leq \lim_{n\to\infty} \frac{c_1 h + h}{c_2 h + h} = \frac{c_1 + 1}{c_2 + 1}$$

Since $c_1 < c_2$, $c_1 + 1 < c_2 + 1$ so $\dfrac{c_1 + 1}{c_2 + 1} < 1$. Thus, $f + h \prec g + h$.

From these results and Corollary 4.3.1, it can trivially be shown that the same statement is true when $\prec$ is replaced with $\preceq$. □